


The Basic Concept of RNN and Google's Inception Model

2016.07.25

서지혜


RNN

► Recurrent Neural Networks




Recurrent Neural Networks have loops.

RNN



An unrolled recurrent neural network.

RNN




$$s_t = \tanh(Ux_t + Ws_{t-1})$$

$$o_t = \text{softmax}(Vs_t)$$


The Problem of Long-Term Dependencies

“the clouds are in the sky”



“I grew up in France...

I speak fluent French.”




LSTM Networks

- ▶ Long Short Term Memory networks (LSTMs)
 - A special kind of RNN
 - Capable of learning long-term dependencies
 - Tremendously well on a large variety of problems
 - Now widely used

LSTM Networks


RNN

a single layer




LSTM

four interacting layers




The Core Idea Behind LSTMs

► The cell state



Gate




remove or add information to the cell state,
carefully regulated by structures called gates


Gates are a way to optionally
let information through

Various types RNN


one to one




one to many



many to one



many to many



many to many




Image Classification Models

- ▶ TensorFlow Image Processing Tutorials
 - CIFAR-10

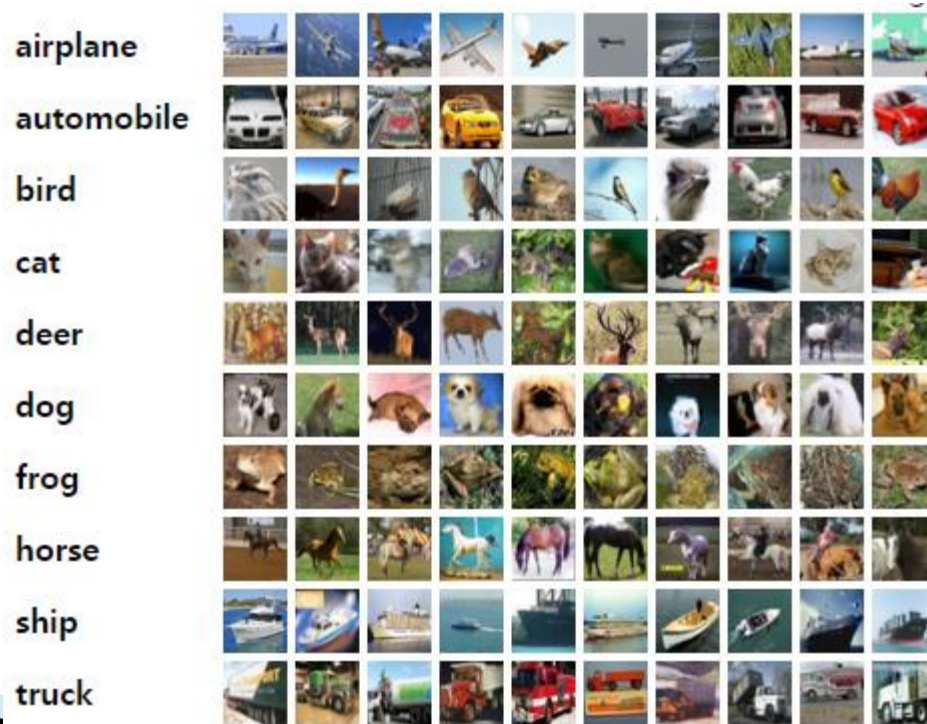


Image Classification Models

- ▶ TensorFlow Image Processing Tutorials
 - ImageNet

```
cd tensorflow/models/image/imagenet  
python classify_image.py
```



```
giant panda, panda, panda bear, coon bear, Ailuropoda melanoleuca (score = 0.88493)  
indri, indris, Indri indri, Indri brevicaudatus (score = 0.00878)  
lesser panda, red panda, panda, bear cat, cat bear, Ailurus fulgens (score = 0.00317)  
custard apple (score = 0.00149)  
earthstar (score = 0.00127)
```


Network

▶ GoogleLeNet

- 1st at ILSVRC 2014
- Inception module

▶ Problems


- Large number of parameters
- Dramatically increased use of computational resources
- Vanishing gradient → overfitting, slow speed



Google : Inception

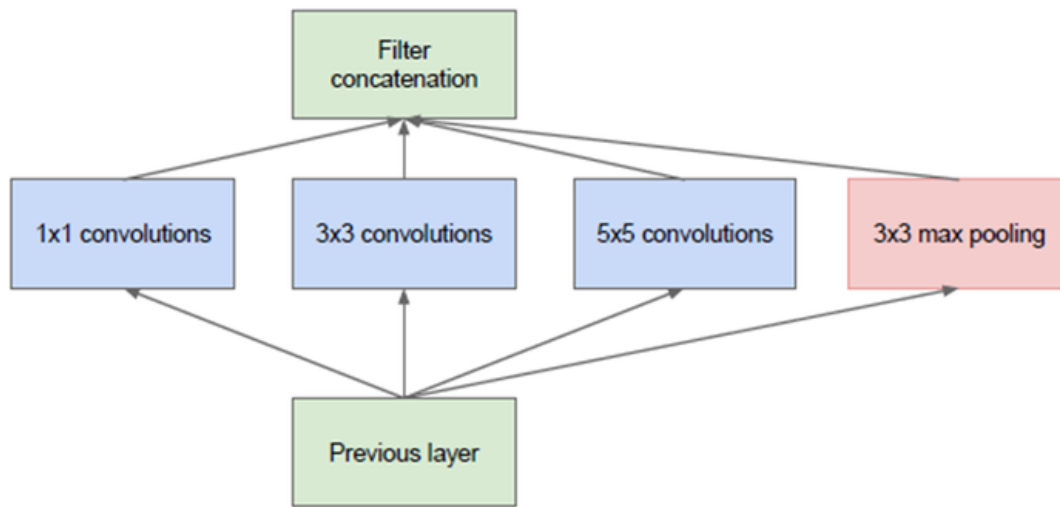


C. Szegedy et al, "Going Deeper with Convolutions" (CVPR 2015)

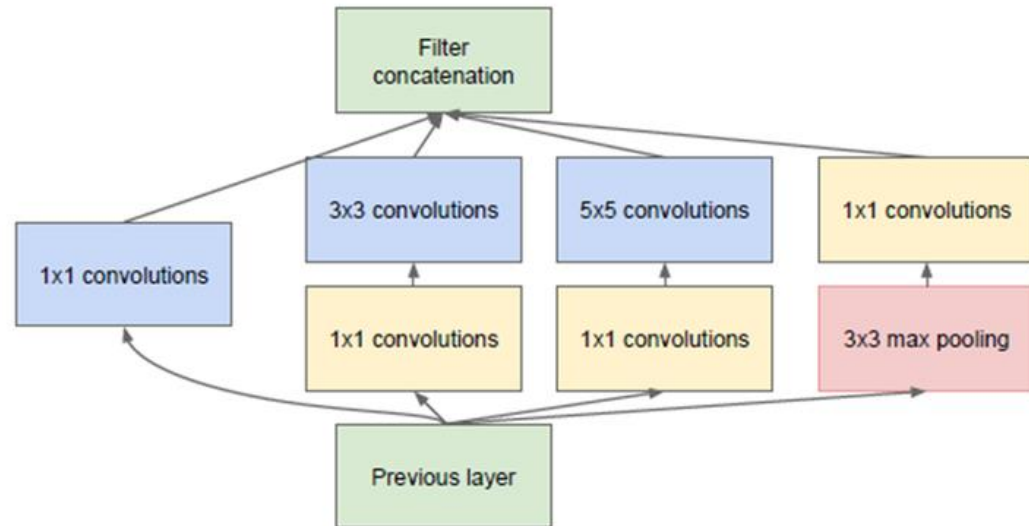


Google : Inception

► Inception Module



(a) Inception module, naïve version



(b) Inception module with dimensionality reduction

Google : Inception

▶ Solution

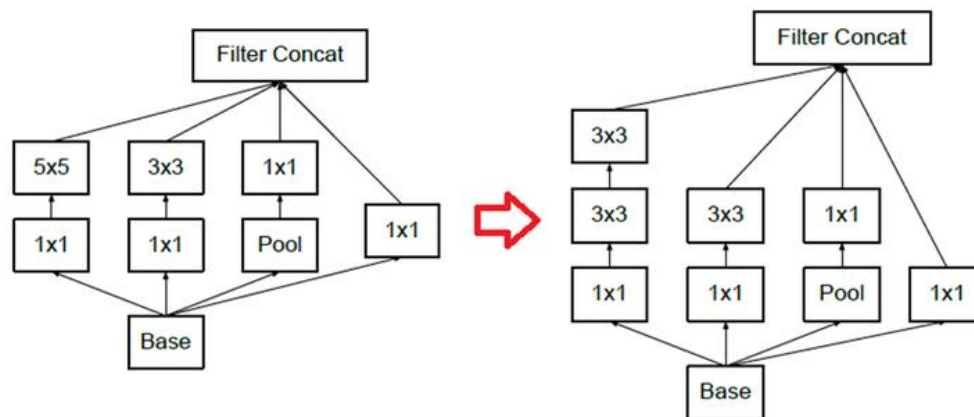
- 1x1 Convolutional Layer

- Auxiliary Classifier

Regularizer 와 같은 역할로 gradient 값을 증가시킨다.

- Factorization into smaller convolutions

큰 conv 커널 → 여러 개의 작은 conv 커널로

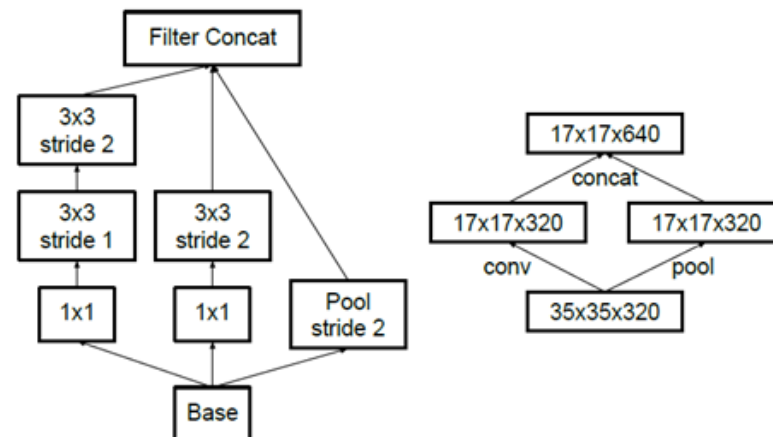


Google : Inception

► Solution

- Grid size reduction

Pooing & conv 나란히 적용




Krizhevsky, Suskever and Hinton (2012)

params *FLOPs*

60M 2B

AlexNet
(8 layers)



Szegedy et al (2014)

5M 1.5B

GoogleLeNet
(22 layers)




Google : Inception

▶ Evaluation result

- Top-1 error : 21.2%
- Top-5 error : 5.6%

▶ The Model Architecture



Train your own image classifier with Inception in TensorFlow

- ▶ Ubuntu 14.04 (64 bit)
- ▶ Python 2.7.6
- ▶ TensorFlow 0.8.0 (CPU only)
- ▶ Bazel 0.2.3

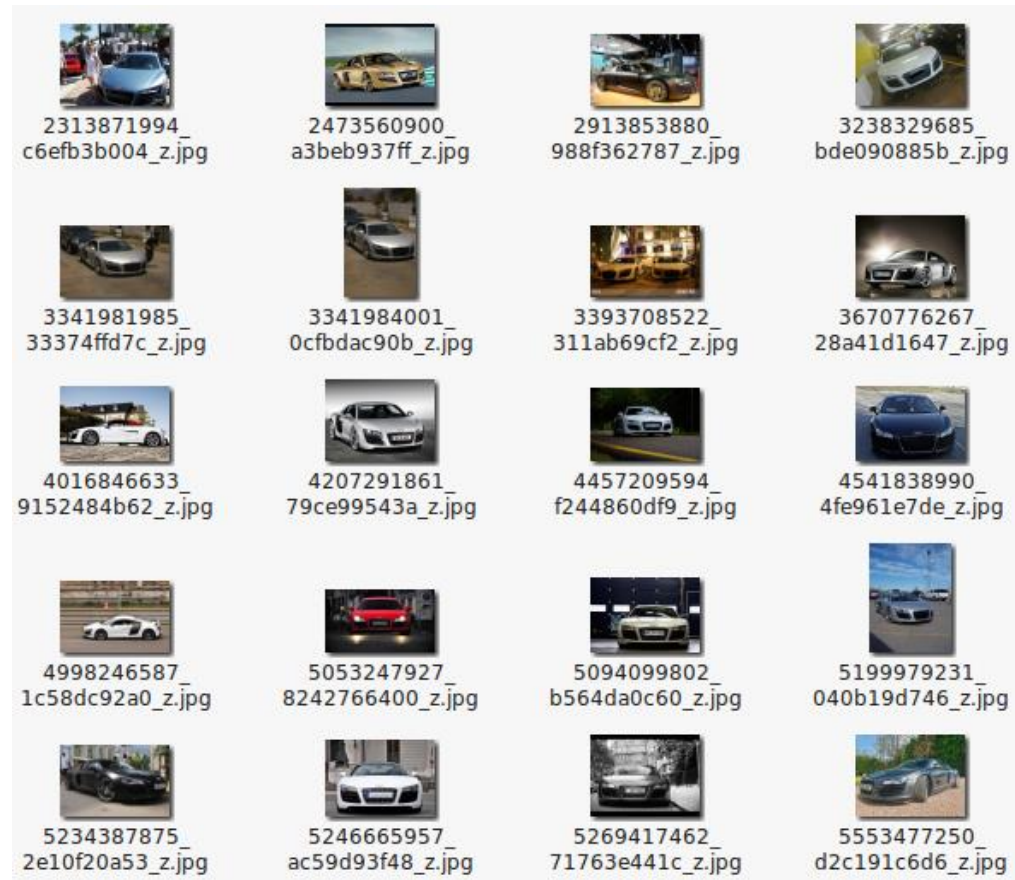



TensorFlow



Train your own image classifier with Inception in TensorFlow

- ▶ Training set – 91.5 MB (854)



Train your own image classifier with Inception in TensorFlow

```
hadoop@hadoop-Samsung-Desktop-System ~/tensorflow $ bazel-bin/tensorflow/examples/image_retrain/retrain --image_dir ~/car_images  
>> Downloading inception-2015-12-05.tgz 26.1%
```

```
hadoop@hadoop-Samsung-Desktop-System ~/tensorflow $ bazel-bin/tensorflow/examples/image_retrain/retrain --image_dir ~/car_images  
Looking for images in 'audi_A1'  
Looking for images in 'bmw_I8'  
Looking for images in 'bmw_X1'  
Looking for images in 'benz_s-class'  
Looking for images in 'hyundai_I40'  
Looking for images in 'hyundai_grandeur'  
Looking for images in 'benz_GLC'  
Looking for images in 'audi_R8'  
Creating bottleneck at /tmp/bottleneck/audi_A1/audi348239052.jpg.txt  
W tensorflow/core/framework/op_def_util.cc:331] Op is deprecated. It will cease to work in GraphDef version 9. Use tf.nn.batch_normalization().  
Creating bottleneck at /tmp/bottleneck/audi_A1/6550727625_89fab191b1_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/9206852839_efc02ec228_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/6548280563_e240dd975b_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/15786813071_acdbc38ba7_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/15603857960_d09b36269a_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/8683157608_817e7a41c1_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/8097947940_09bd8d58ac_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/6548281961_f803d7d4fa_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/12149349514_2f7e7871a4_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/audi374289472.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/6548280721_f04003f043_z.jpg.txt  
Creating bottleneck at /tmp/bottleneck/audi_A1/4749763409_1d3050f247_z.jpg.txt
```

Train your own image classifier with Inception in TensorFlow

```
2016-06-07 11:03:32.379385: Step 0: Train accuracy = 26.0%
2016-06-07 11:03:32.379492: Step 0: Cross entropy = 2.042176
2016-06-07 11:03:32.430124: Step 0: Validation accuracy = 18.0%
2016-06-07 11:03:32.930563: Step 10: Train accuracy = 46.0%
2016-06-07 11:03:32.930648: Step 10: Cross entropy = 1.926688
2016-06-07 11:03:32.977640: Step 10: Validation accuracy = 36.0%
2016-06-07 11:03:33.479548: Step 20: Train accuracy = 59.0%
2016-06-07 11:03:33.479630: Step 20: Cross entropy = 1.861638
2016-06-07 11:03:33.531690: Step 20: Validation accuracy = 51.0%
```

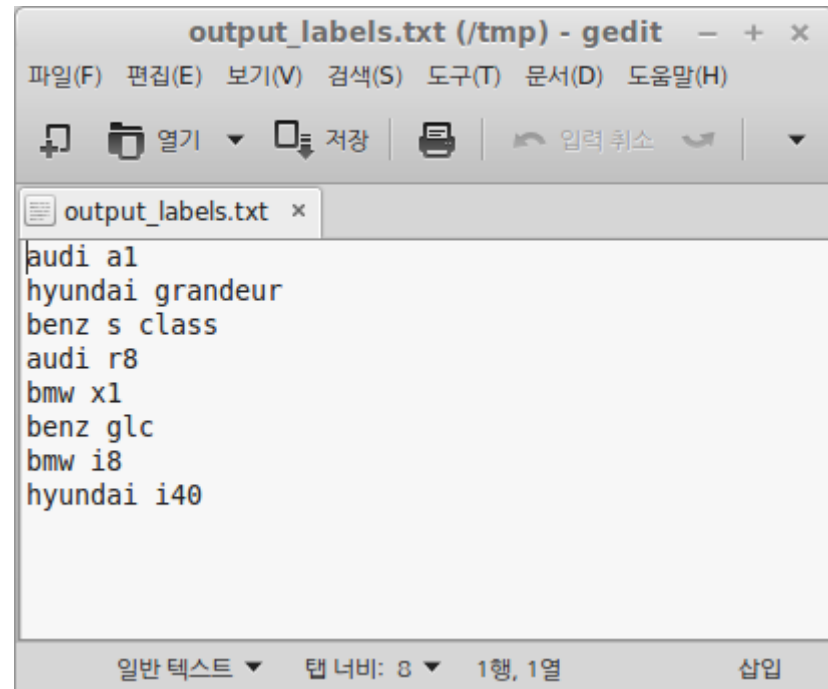

```
2016-06-07 11:05:07.406752: Step 1720: Train accuracy = 92.0%
2016-06-07 11:05:07.406836: Step 1720: Cross entropy = 0.458142
2016-06-07 11:05:07.454632: Step 1720: Validation accuracy = 63.0%
2016-06-07 11:05:07.971826: Step 1730: Train accuracy = 93.0%
2016-06-07 11:05:07.971910: Step 1730: Cross entropy = 0.404417
2016-06-07 11:05:08.020165: Step 1730: Validation accuracy = 67.0%
```

```
2016-06-07 11:05:27.839287: Step 2090: Train accuracy = 93.0%
2016-06-07 11:05:27.839367: Step 2090: Cross entropy = 0.370275
2016-06-07 11:05:27.886765: Step 2090: Validation accuracy = 73.0%
2016-06-07 11:05:28.397926: Step 2100: Train accuracy = 95.0%
2016-06-07 11:05:28.397996: Step 2100: Cross entropy = 0.328979
2016-06-07 11:05:28.447660: Step 2100: Validation accuracy = 70.0%
```

```
2016-06-07 11:07:12.770920: Step 3990: Train accuracy = 98.0%
2016-06-07 11:07:12.771004: Step 3990: Cross entropy = 0.231101
2016-06-07 11:07:12.816862: Step 3990: Validation accuracy = 64.0%
2016-06-07 11:07:13.277159: Step 3999: Train accuracy = 99.0%
2016-06-07 11:07:13.277288: Step 3999: Cross entropy = 0.209968
2016-06-07 11:07:13.329315: Step 3999: Validation accuracy = 65.0%
Final test accuracy = 78.8%
Converted 2 variables to const ops.
```


Train your own image classifier with Inception in TensorFlow

output_labels


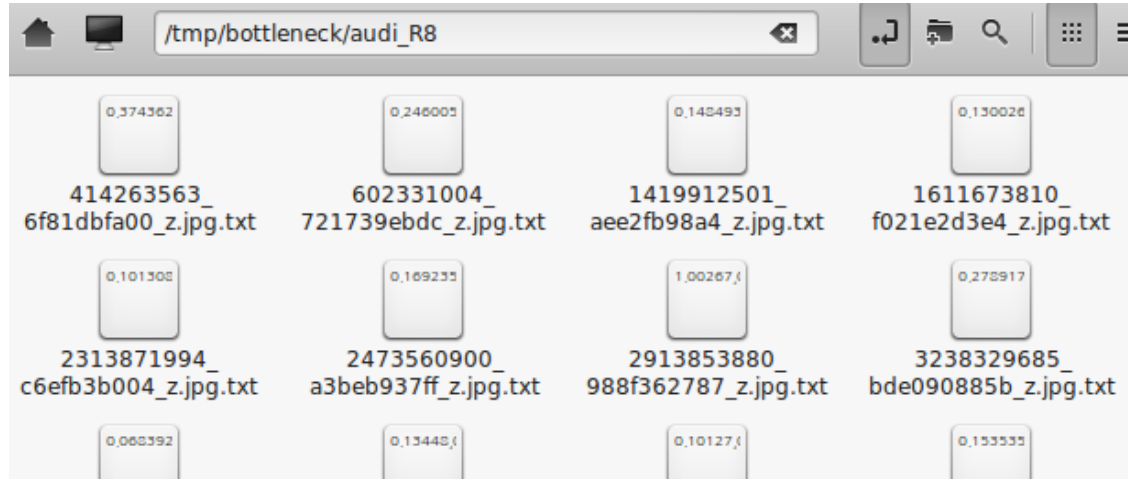



A screenshot of a gedit text editor window titled 'output_labels.txt (/tmp) - gedit'. The window shows the contents of the file 'output_labels.txt' in a monospaced font. The text is as follows:

```
audi a1  
hyundai grandeur  
benz s class  
audi r8  
bmw x1  
benz glc  
bmw i8  
hyundai i40
```

The editor's status bar at the bottom indicates '일반 텍스트' (Plain Text), '탭 너비: 8' (Tab width: 8), and '1행, 1열' (1 line, 1 column).

Train your own image classifier with Inception in TensorFlow



Train your own image classifier with Inception in TensorFlow

Audi a1



```
audi a1 (0): 0.808858  
bmw x1 (4): 0.103799  
hyundai grandeur (1): 0.0264274  
bmw i8 (6): 0.0263073  
benz glc (5): 0.013007
```



```
audi a1 (0): 0.99349  
benz glc (5): 0.00309179  
benz s class (2): 0.00120927  
hyundai grandeur (1): 0.00116204  
bmw x1 (4): 0.000675393
```



```
audi a1 (0): 0.758456  
benz s class (2): 0.121226  
hyundai grandeur (1): 0.0917348  
bmw x1 (4): 0.0209011  
benz glc (5): 0.00401876
```

Bmw x1



```
benz glc (5): 0.497991  
audi a1 (0): 0.248487  
bmw x1 (4): 0.197361  
hyundai i40 (7): 0.0242653  
benz s class (2): 0.0195948
```



```
bmw x1 (4): 0.876124  
audi a1 (0): 0.0725529  
benz glc (5): 0.02024  
benz s class (2): 0.0144952  
hyundai i40 (7): 0.0097391
```



```
bmw x1 (4): 0.650448  
benz s class (2): 0.160187  
audi a1 (0): 0.111266  
benz glc (5): 0.0699459  
hyundai grandeur (1): 0.00753701
```


Train your own image classifier with Inception in TensorFlow

Audi r8



✓
hyundai i40 (7): 0.556305
audi r8 (3): 0.145987
hyundai grandeur (1): 0.0971609
audi a1 (0): 0.0867918
bmw i8 (6): 0.0572504



audi r8 (3): 0.908869
audi a1 (0): 0.0571011
bmw i8 (6): 0.0161388
hyundai grandeur (1): 0.0137875
bmw x1 (4): 0.00194827



✓
bmw i8 (6): 0.719148
audi a1 (0): 0.159573
audi r8 (3): 0.0759141
hyundai grandeur (1): 0.0228197
hyundai i40 (7): 0.0165728

Hyundai i40



hyundai i40 (7): 0.880916
bmw i8 (6): 0.0561588
hyundai grandeur (1): 0.032098
audi a1 (0): 0.0166874
bmw x1 (4): 0.0114884



✓
hyundai grandeur (1): 0.640598
audi a1 (0): 0.186405
hyundai i40 (7): 0.0819051
bmw i8 (6): 0.0256575
benz s class (2): 0.0240114



hyundai i40 (7): 0.91717
audi a1 (0): 0.0275488
hyundai grandeur (1): 0.0229953
benz glc (5): 0.0159541
benz s class (2): 0.0120241

Train your own image classifier with Inception in TensorFlow

Training set

91.5 MB (854)  97.8 MB (902)

Final test accuracy

78.8 %  77.4 %

Train your own image classifier with Inception in TensorFlow

Bmw x1



```
benz glc (5): 0.497991  
audi a1 (0): 0.248487  
bmw x1 (4): 0.197361  
hyundai i40 (7): 0.0242653  
benz s class (2): 0.0195948
```



```
benz glc (5): 0.438972  
audi a1 (0): 0.291268  
bmw x1 (4): 0.214732  
hyundai i40 (7): 0.0194838  
benz s class (2): 0.0177405
```

Audi r8



```
bmw i8 (6): 0.719148  
audi a1 (0): 0.159573  
audi r8 (3): 0.0759141  
hyundai grandeur (1): 0.0228197  
hyundai i40 (7): 0.0165728
```



```
bmw i8 (6): 0.725414  
audi a1 (0): 0.11334  
audi r8 (3): 0.10061  
hyundai i40 (7): 0.0346341  
hyundai grandeur (1): 0.0214064
```

Hyundai i40



```
hyundai grandeur (1): 0.640598  
audi a1 (0): 0.186405  
hyundai i40 (7): 0.0819051  
bmw i8 (6): 0.0256575  
benz s class (2): 0.0240114
```



```
hyundai i40 (7): 0.539358  
hyundai grandeur (1): 0.260657  
audi a1 (0): 0.137662  
bmw i8 (6): 0.0194107  
audi r8 (3): 0.0155325
```

References

- ▶ C. Szegedy, W. Liu, Y. Jia, P. Sermanet, S. Reed, D. Anguelov, D. Erhan, V. Vanhoucke, and A. Rabinovich. Going deeper with convolutions . In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2015.
- ▶ C. Szegedy, V. Vanhoucke, S. Ioffe, J. Shlens, and Z. Wojna. Rethinking the inception architecture for computer vision. arXiv preprint arXiv:1512.00567, 2015.
- ▶ TensorFlow official home page (<https://www.tensorflow.org>)
- ▶ TensorFlow Github
(<https://github.com/tensorflow/models/tree/master/inception>)
- ▶ Laonpeople Naver Blog (<http://laonple.blog.me>)