

A review of Random feedback weights support learning in deep neural networks

SangHyun Yi

sangyi92@snu.ac.kr

November 7, 2016

Table of contents

1 The idea of Random feedback

- Recap backprop
- Backprop is not the way a brain learns
- Random feedback alignment

2 Empirical results of random feedback

- Linear function-approximation
- Nonlinear, real-world problems

3 Why it works? Theorems

- Why it works?
- Condition for alignment to reduce error to zero
- Gauss-Newton modification of backprop
- B acts like the pseudoinverse of W

The idea of Random feedback

Recap backprop




Figure: A sketch of the backprop


Recap backprop

n layer MLP

$\sigma(\mathbf{x})$: *activation function*

$$\mathbf{y}_k = \sigma(\mathbf{x}_k)$$

W_m : *a weight multiplied to m th layer to m + 1st layer*



$$\delta_n = \frac{\partial \mathcal{L}}{\partial \mathbf{x}_n} = \frac{\partial \mathbf{y}_n}{\partial \mathbf{x}_n} \frac{\partial \mathcal{L}}{\partial \mathbf{y}_n} = \text{diag}(\sigma'(\mathbf{x}_n)) \frac{\partial \mathcal{L}}{\partial \mathbf{y}_n}$$

$$\delta_{k-1} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}_{k-1}} = \frac{\partial \mathbf{x}_k}{\partial \mathbf{x}_{k-1}} \frac{\partial \mathcal{L}}{\partial \mathbf{x}_k} = \text{diag}(\sigma'(\mathbf{x}_{k-1})) W_{k-1}^t \delta_k$$

$$\Delta W_k = \frac{\partial \mathcal{L}}{\partial W_k} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}_{k+1}} \frac{\partial \mathbf{x}_{k+1}}{\partial W_k} = \delta_{k+1} \mathbf{y}_k^t$$

$$W_k^{(t+1)} = W_k^{(t)} - \eta \Delta W_k^{(t)}$$

Backprop is not the way a brain learns



- Backprop requires a precise transport of synaptic weight information. i.e. impossible
- Synapses in brain functions in a unidirectional way. Not bidirectional.
- Reinforcement learning? Shallow mechanism?
- How about putting a random matrix B to the place of W^t ?


Random feedback alignment

n layer MLP

$\sigma(\mathbf{x})$: *activation function*

$$\mathbf{y}_k = \sigma(\mathbf{x}_k)$$

W_m : *a weight multiplied to m th layer to m + 1st layer*



$$\delta_n = \frac{\partial \mathcal{L}}{\partial \mathbf{x}_n} = \frac{\partial \mathbf{y}_n}{\partial \mathbf{x}_n} \frac{\partial \mathcal{L}}{\partial \mathbf{y}_n} = \text{diag}(\sigma'(\mathbf{x}_n)) \frac{\partial \mathcal{L}}{\partial \mathbf{y}_n}$$

$$\delta_{k-1} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}_{k-1}} = \frac{\partial \mathbf{x}_k}{\partial \mathbf{x}_{k-1}} \frac{\partial \mathcal{L}}{\partial \mathbf{x}_k} = \text{diag}(\sigma'(\mathbf{x}_{k-1})) B_{k-1} \delta_k$$

$$\Delta W_k = \frac{\partial \mathcal{L}}{\partial W_k} = \frac{\partial \mathcal{L}}{\partial \mathbf{x}_{k+1}} \frac{\partial \mathbf{x}_{k+1}}{\partial W_k} = \delta_{k+1} \mathbf{y}_k^t$$

$$W_k^{(t+1)} = W_k^{(t)} - \eta \Delta W_k^{(t)}$$

Empirical results of random feedback


Linear function-approximation

T : target linear function

$$\mathbf{y} = W\mathbf{h}$$

$$\mathbf{h} = W_0\mathbf{x}$$

$$\mathcal{L} = \frac{1}{2}\mathbf{e}^t\mathbf{e}, \mathbf{e} = \mathbf{y}^* - \mathbf{y}, \mathbf{y}^* = T\mathbf{x}$$




Linear function-approximation

$$\Delta W_0 \propto \frac{\partial \mathcal{L}}{\partial W_0} = \frac{\partial \mathcal{L}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial W_0} = -(W^t \mathbf{e}) \mathbf{x}^t$$

$$\Delta W \propto \frac{\partial \mathcal{L}}{\partial W} = -\mathbf{e} \mathbf{h}^t$$

$$\Delta \mathbf{h}_{BP} = W^t \mathbf{e}$$

$$\Delta \mathbf{h}_{FA} = B \mathbf{e}$$



Linear function-approximation

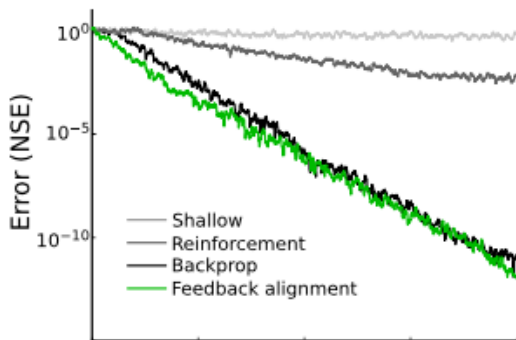


Figure: Four algorithms learn to mimic a linear function

- $NSE(\mathbf{y}^*, \mathbf{y}) = \frac{MSE(\mathbf{y}^*, \mathbf{y})}{MSE(\mathbf{y}^*, 0)}$

Linear function-approximation




Figure: Angle between the hidden-unit update vector prescribed by feedback alignment and that prescribed by backprop

- $\mathbf{e}^t W B \mathbf{e} > 0$
- This alignment of the Δh 's implies that B has begun to act like W^t .

Nonlinear, real-world problems

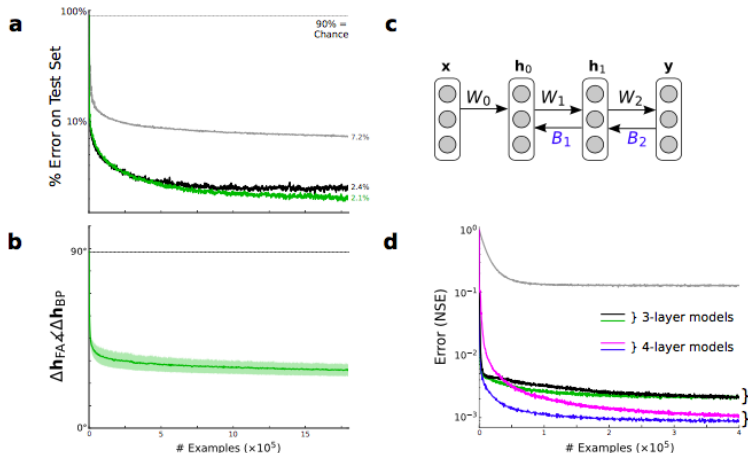


Figure: **a**: MNIST, **b**: angle between FA and BP, **c**: FA in deep network, **d**: nonlinear function-approximation

Nonlinear, real-world problems




Figure: Receptive fields for 100 randomly selected hidden units

Why it works? Theorems

Why it works?




Figure: Vector flow field (small arrows) demonstrates the evolution of W_0 and W during feedback alignment

- Linear function
- $T = 1$ and B is 'randomly' chosen to be 1
- Thick lines are solutions. i.e. $T = WW_0 = 1$
- grey: $e^t W B e > 0$, black: $e^t W B e < 0$, dash: unstable solution

Why it works?




Figure: Vector flow field (small arrows) demonstrates the evolution of W_0 and W during feedback alignment

- Upper right hyperbola is a set of solutions where $W > 0$
- $\mathbf{e}^t W B \mathbf{e} > 0 \forall e$
- Meaning W has evolved so that the feedback matrix B is delivering useful teaching signals.

Condition for alignment to reduce error to zero

$$\mathbf{y} = W\mathbf{h}$$

$$\mathbf{h} = A\mathbf{x}$$

$$E := T - WA$$

$$\mathbf{e} = E\mathbf{x}$$

$$\mathcal{L} = \frac{1}{2}\mathbf{e}^t\mathbf{e}, \mathbf{e} = \mathbf{y}^* - \mathbf{y}, \mathbf{y}^* = T\mathbf{x}$$

$$\Delta W = \eta \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial W} = \eta E \mathbf{x} \mathbf{x}^t A^t$$

$$\Delta A = \eta \frac{\partial \mathbf{y}}{\partial \mathbf{h}} \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{h}}{\partial A} = \eta B E \mathbf{x} \mathbf{x}^t$$

Condition for alignment to reduce error to zero

Assume $\mathbf{x} \sim iid \mathbf{N}(0, I)$.

$$\Delta W = \eta[E\mathbf{x}\mathbf{x}^t A^t] = \eta E A^t$$

$$\Delta A = \eta[B E\mathbf{x}\mathbf{x}^t] = \eta B E$$

$$\dot{W} = E A^t$$

$$\dot{A} = B E$$

Then by integration,

$$B W + W^t B^t = A A^t + C$$

Condition for alignment to reduce error to zero

Theorem (1)

Given the learning dynamics

$$\dot{W} = EA^t$$

$$\dot{A} = BE$$

and assuming that the constant C in the equation is zero and that the matrix B satisfies

$$B^+B = I$$

then

$$\lim_{t \rightarrow \infty} E = 0$$

Condition for alignment to reduce error to zero

Sketch of the proof

- Let $V := \text{tr}(BEE^t B^t)$
- V is lower bounded, \dot{V} is negative semi-definite, and \dot{V} is uniformly continuous in time, which is satisfied since \ddot{V} is finite.
- $\therefore \dot{V} \rightarrow 0$ as $t \rightarrow \infty$, which means V converges to some value.
- Since $\dot{V} = -2\text{tr}(BEA^t AE^t B^t) - \text{tr}(A^t BEE^t B^t A) = 0$,
 $\text{tr}(BEA^t AE^t B^t) = 0$ i.e. $BEA^t = 0$
- $\therefore B^+ BEA^t = EA^t = 0$
- Since $\dot{W} = EA^t$, W is constant.
- $\therefore AA^t = BW + W^t B^t$ is constant.
- Since $0 = BEA^t = BTA^t - BWAA^t$, BTA^t is constant.
- $\frac{dBTA^t}{dt} = 0 = BTE^t B^t$
- $0 = BTE^t B^t = TE^t = ET^t$
- $\therefore EE^t = ET^t - EA^t W^t = 0$

Gauss-Newton modification of backprop

$$\mathbf{y} = W\mathbf{h}$$

$$\mathbf{h} = \sigma(A\mathbf{x})$$

$$E := T - WA$$

$$\mathbf{e} = E\mathbf{x}$$

$$\mathcal{L} = \frac{1}{2}\mathbf{e}^t\mathbf{e}, \mathbf{e} = \mathbf{y}^* - \mathbf{y}, \mathbf{y}^* = T\mathbf{x}$$

$$\mathcal{L}_x = \mathbf{e}_x^t\mathbf{e}, \mathcal{L}_{xx} = \mathbf{e}_{xx}^t\mathbf{e} + \mathbf{e}_x^t\mathbf{e}_x \approx \mathbf{e}_x^t\mathbf{e}_x \text{ when } \mathbf{e} \text{ is small.}$$

$$\Delta x = -\mathcal{L}_{xx}^{-1}\mathcal{L}_x^t \approx -(\mathbf{e}_x^t\mathbf{e}_x)^{-1}\mathbf{e}_x^t\mathbf{e} = \mathbf{e}_x^+\mathbf{e}$$

Theorem (2)

Using W^+ instead of W^t in backprop algorithm makes the algorithm behave like Gauss-Newton method.

Sketch of the proof

- $\Delta \mathbf{h}_{GN} = -\mathbf{e}_h^+ \mathbf{e} = -W^+ \mathbf{e}$
- $\Delta A_{BP} = -\eta \text{Diag}(\sigma'(\mathbf{x})) W^t \mathbf{e} \mathbf{x}^t$
- $\Delta A_{PBP} = -\eta \text{Diag}(\sigma'(\mathbf{x})) W^+ \mathbf{e} \mathbf{x}^t$
- $\Delta \mathbf{h}_{PBP} = \Delta(\sigma(A\mathbf{x})) \approx \text{Diag}(\sigma'(A\mathbf{x})) \Delta(A\mathbf{x}) = \text{Diag}(\sigma'(A\mathbf{x})) \Delta A_{PBP} \mathbf{x} = (\text{Diag}(\sigma'(A\mathbf{x})))^2 \Delta \mathbf{h}_{GN} \mathbf{x}^t \mathbf{x}$
- $\therefore \Delta \mathbf{h}_{PBP} \propto \Delta \mathbf{h}_{GN}$ elementwise.

B acts like the pseudoinverse of W

$$\mathbf{y} = W\mathbf{h}$$

$$\mathbf{h} = A\mathbf{x}$$

$$E := T - WA$$

$$\mathbf{e} = E\mathbf{x}$$

$$\mathcal{L} = \frac{1}{2}\mathbf{e}^t\mathbf{e}, \mathbf{e} = \mathbf{y}^* - \mathbf{y}, \mathbf{y}^* : \text{True value}$$

$$W^{(t+1)} = W^{(t)} + \Delta W^{(t)}$$

$$A^{(t+1)} = A^{(t)} + \Delta A^{(t)}$$

B acts like the pseudoinverse of W

$$\Delta W = \eta_W \mathbf{e} \mathbf{h}^t$$

$$\Delta A = \eta_A B \mathbf{e} \mathbf{x}^t$$

$$\mathbf{h}^{(t+1)} = A^{(t+1)} \mathbf{x} = (A^{(t)} + \Delta A^{(t)}) \mathbf{x} = \mathbf{h}^{(t)} + \eta_A B \mathbf{e} \mathbf{x}^t \mathbf{x} = \mathbf{h}^{(t)} + \eta_h B \mathbf{e}$$

Let $\eta_h = \eta_W = \eta$. Then,
 $W^{(t+1)} = W^{(t)} + \Delta W^{(t)}$

$$\mathbf{h}^{(t+1)} = \mathbf{h}^{(t)} + \Delta \mathbf{h}^{(t)}$$

with,

$$\Delta W = \eta \mathbf{e} \mathbf{h}^t$$

$$\Delta \mathbf{h} = \eta B \mathbf{e}$$

B acts like the pseudoinverse of W

Theorem (3)

When W and A are initialized to zero, hidden unit updates prescribed by the forward alignment algorithm, $\Delta \mathbf{h}_{FA}$, are always a positive scalar multiple of the hidden unit updates prescribed by the pseudobackprop algorithm, $\Delta \mathbf{h}_{PBP}$.

B acts like the pseudoinverse of W

Sketch of the proof

- $\exists s_h, s_W$ s.t. $\mathbf{h} = s_h B \mathbf{y}^*$, $W = s_W \mathbf{y}^* (B \mathbf{y}^*)^t$
- Trivial when $\mathbf{h} = 0, W = 0$. Then by induction, above equations hold for every time step.
- $\mathbf{e} = (1 - s_y) \mathbf{y}^*$ with $(1 - s_y)$ a positive scalar.
- $\therefore \Delta \mathbf{h}_{FA} = \eta B \mathbf{e} = \eta (1 - s_y) B \mathbf{y}^*$ and $\Delta \mathbf{h}_{BBP} = \eta W^+ \mathbf{e} = \eta (1 - s_y) W^+ \mathbf{y}^*$
- It suffices to show that $s B \mathbf{y}^* = (\mathbf{y}^* (B \mathbf{y}^*)^t)^+ \mathbf{y}^*$.
- $(\mathbf{y}^* (B \mathbf{y}^*)^t)^+ \mathbf{y}^* = (B \mathbf{y}^*)^{t+} \mathbf{y}^{*+} \mathbf{y}^* = (B \mathbf{y}^*)^{t+} = s B \mathbf{y}^*$ where $s = ((B \mathbf{y}^*)^t (B \mathbf{y}^*))^{-1}$.
- $\therefore \Delta \mathbf{h}_{FA} = s \Delta \mathbf{h}_{BBP}$

B acts like the pseudoinverse of W

